Development accuracy versus cumulative online secure runtime (seed42)

CIFAR-10

CIFAR-100

Tiny-ImageNet

TFE / ABY3

SecretFlow-SPU

Development validation accuracy

70%
60%
50%
40%
30%

70.0%
60.0%
50.0%
40.0%
30.0%

10^2
10^3
10^4
3×10^3
4×10^3
6×10^3
2×10^4

Full Data (n=1)
True Ours (n=1)

Cumulative online secure runtime (seconds, log scale)

Backend-specific measured online boundaries, not a strict hardware-normalized speedup comparison. ε is conditional Stage C only. SPU n=1, no variance shading.